



Open-weight (Qwen, Llama, Mistral) Closed-API zero-shot (GPT-4o, Sonnet 4.5)
Safety + behavior across 6 model families on identical guarded evidence (n=114)


Parse OK rate



Cross-trial leak rate (all 0)



Commit rate



Abstain rate

