

■ Open-weight (Qwen, Llama, Mistral) ■ Closed-API zero-shot (GPT-4o, Sonnet 4.5)  
Safety + behavior across 6 model families on identical guarded evidence (n=114)

